

Job Title: Applied AI Engineer

Location: Pune, India (Hybrid)

Employment Type: Full-time

#### About the Role

We are looking for an Applied AI Engineer to join our product team and build production-grade LLM systems — not demos. You'll design RAG pipelines, agent workflows, and APIs that real users depend on daily.

#### What You'll Do

- Build and ship RAG pipelines and multi-agent workflows using LangChain, LangGraph, or similar frameworks
- Develop FastAPI backends with streaming (SSE), async processing, and solid error handling
- Design model routing, fallbacks, and validation layers to reduce hallucinations in production
- Work with PostgreSQL, vector stores (FAISS, Qdrant), and Redis for retrieval and caching
- Deploy and monitor AI services on AWS or Azure (SageMaker, Azure OpenAI, etc.)
- Collaborate with product and engineering on NL→SQL, document intelligence, or similar use cases

#### Required Qualifications

- 1–3 years of professional experience building AI/ML systems (strong portfolio counts)
- Hands-on experience with Python, FastAPI, and at least one LLM framework (LangChain, LangGraph, or Vercel AI SDK)
- Production experience with RAG, prompt design, and LLM evaluation
- Familiarity with PostgreSQL and at least one vector database
- Comfort with Docker, CI/CD, and cloud deployment (AWS or Azure)

#### Nice to Have

- LoRA/QLoRA fine-tuning experience
- Multi-model routing and cost optimization
- Document parsing / PDF intelligence pipelines
- MLOps tooling (MLflow, Weights & Biases)
- Experience leading small engineering teams or presenting to non-technical stakeholders

#### What We Offer

- Competitive salary and equity
- Hybrid work in Pune
- Remote-friendly Fridays
- Learning budget for conferences and courses

#### How to Apply

Submit your resume and a short note on a production AI system you've built. We are an equal opportunity employer.